A Probabilistic Model for High Recall Feature Selection

Paper # 1253

Abstract

Feature selection is important in many practical supervised learning settings to address computational constraints or to prevent overfitting when the data may be high-dimensional but contain relatively few samples. While a wide variety of feature selection methods have been proposed in the literature, little research seems to focus on feature selection specifically targeted to improve recall in the case of binary classification. In this paper we propose a novel probabilistic voting model of feature selection and argue that choosing features so as to maximize likelihood objectives w.r.t. this model provides an effective method for high recall feature selection. For each objective, we derive an efficient formula for feature selection in the filtering framework (i.e., greedy forward selection) and we empirically compare the resulting feature selection criteria to a wide variety of existing methods. Results show that our high recall feature selection approach does indeed improve recall with improvements noticeable when there is a high feature to data ratio. Such results provide a novel and efficient feature selection algorithm to target high-recall classification tasks.

Introduction

Feature selection is important in many practical supervised learning settings to address computational constraints (e.g., in large-scale or online learning) or to prevent overfitting when the data may be high-dimensional but contain relatively few samples (e.g., as may occur in medical or bioinformatics domains where features are abundant but data is costly to obtain) Guyon and Elisseeff (2003).

While a wide variety of feature selection methods have been proposed in the literature, little research seems to focus on feature selection specifically targeted to improve recall (minimization of false negatives) in binary classification. One reason for this is simply that most feature selection methods are agnostic to the particular supervised learning task – applying to tasks from classification to regression – and hence are not focused on performance properties specific to binary classifiers. Yet, recall is an important aspect of many binary classification problems (e.g., minimizing false

negatives in the identification of cancerous tumors) and it is critical to have feature selection algorithms that are targeted to maintain high recall.

Naturally though, it would not make sense for a classifier or feature selection method to focus on recall alone since the classifier which always predicts *true* (independent of the data) obtains optimal recall performance. To trade off precision with recall in order to maintain high accuracy, one often uses a geometric average of the two instead (known as F-score) but in practice most classifiers directly optimize accuracy or some surrogate, e.g., a convex surrogate of 0-1 loss as in the SVM or maximum (conditional) likelihood as in Naive Bayes or logistic regression. So how can one achieve a high recall SVM, Naive Bayes, or logistic regression classifier that already has a well-defined accuracy-focused optimization criterion? The idea we pursue in this paper is that feature selection can help modulate the recall performance of existing classifiers by encouraging selection of features that cover more of the true cases in the data (hence discouraging false negatives).

In this paper we propose a novel probabilistic voting model of feature selection with the intent of encouraging false negative reduction while still focusing on accuracy. We argue that choosing features so as to maximize likelihood objectives w.r.t. this model provides an effective method for high recall feature selection. Specifically, for each objective, we derive an efficient formula for feature selection in the filtering framework (i.e., greedy forward selection) and we empirically compare the resulting feature selection criteria to a wide variety of existing methods.

Our results demonstrate that our high recall feature selection approach does indeed improve recall with improvements noticeable when there is a high feature to data ratio. Such results provide a novel and efficient feature selection algorithm to target high-recall classification tasks.

Classification and Feature Selection

In this section, we briefly define the task of binary classification along with standard definitions of performance metrics we may wish to optimize. We then follow this by a discussion of feature selection and existing criteria proposed in the literature.

In the binary classification task, we assume we are given data $D = \{(\vec{x}^d, y^d)\}$ consisting of pairs of real-valued raw

input vectors $\vec{x}^d \in \mathbb{R}^n$ of length n (e.g., the results of n different medical tests) and *actual* binary class label $y^d \in \{0(\text{false}), 1(\text{true})\}$ (where we often write F for false and T for true). A binary classifier is a function $C : \vec{x}^d \rightarrow y^d$ such that given a new unlabeled raw feature vector, $C(\vec{x}^d)$ produces a *predicted* classification.

Given a trained classifier C and a dataset D, we can build the well-known contingency table

	Actual T	Actual F
Predicted T	TP	FP
Predicted F	FN	TN

where the four entries represent the counts of true positives (TP), false positives (FP), false negatives (FN) and true negatives (TN) and sum to the total amount of data (i.e., $TP + FP + FN + TN = |D|$). Each table entry represents the count of data for which the respective row matched the predicted classification $C(\vec{x}^d)$ and the respective column matched the actual label y^d. Given these definitions, we can easily define

$$\text{Accuracy} = \frac{TP + TN}{TP + FP + FN + TN}$$
$$\text{Recall} = \frac{TP}{TP + FN}$$

where accuracy represents the overall fraction of correct classification, but recall represents the overall fraction of true labeled data that has been classified as true (i.e., recalled). For somewhat rare events such as medical diagnosis of cancer, we certainly care about accuracy, but we may also want to place additional emphasis on recall performance so as to avoid the occurrence of false negatives (cases of cancer that were missed by the classifier). Of course false positives are also a problem, but additional tests would rule these out and hence not as critical of a classification failure as missing a potential cancer diagnosis.

Previously we did not specify exactly how $C(\vec{x}^d)$ learns from the raw data $\vec{x}^d$. In general, practitioners often process the raw input vectors $\vec{x}^d$ into a subset of features that we'll denote $f_1, \ldots, f_k$. Whereas the raw input may consist of the results of individual tests, a feature $f_i \in \mathbb{R}$ may represent some nonlinear function of one or more tests deemed to be relevant to the classification task. Hence we might more appropriately write a classifier as $C(\vec{x}^d, f_1, \ldots, f_k)$ to represent the raw input data and the features of the data that the classifier may use. In the case of high-dimensional data or few data samples, we may wish to limit the set of features generated to improve classifier performance and this is the task of feature selection — select $\{f_1, \ldots, f_k\}$ to optimize performance. There are already a variety of feature selection methods that we outline next.

Existing Feature Selection Algorithms

Existing feature selection algorithms fall into categories of filter methods and wrapper methods Guyon and Elisseeff (2003). While wrapper methods, such as SVM-RFE Guyon et al. (2002), select features by evaluating their usefulness to a given classifier, filter methods evaluate feature subsets according to certain properties of themselves. Within filter methods, feature selection algorithms can be further classified into ranking methods and subset methods. The former evaluates each features independently, while the latter evaluates a subset at a time Brown et al. (2012). Listed below is a brief introduction to algorithms to which we compared our proposed model.

Correlation based rank is a naïve ranking algorithm that ranks features according to their linear correlation with the output, i.e. Pearson's r. More sophisticated ranking methods include:

1. The conditional entropy of class Y given feature X, i.e. $H(Y|X)$, quantifies the amount of information in Y that is not provided by X. Subtracting $H(Y|X)$ from the entropy of class Y, the information gain of class Y given feature X is the mutual information between feature X and class Y, i.e. $I(X;Y)$.

2. The gain ratio of feature X is defined as the information gain of feature X normalized against the entropy of itself, i.e. $\frac{I(X;Y)}{H(X)}$.

3. The symmetric uncertainty between feature X and class Y measures the amount of redundancy between them. It is defined as $U(X,Y) = 2\frac{I(X;Y)}{H(X)+H(Y)}$.

4. The Relief method Kira and Rendell (1992) evaluates the relevance of features to the output class according to how well their values distinguish between nearest instances of the same and different classes.

Correlation based subset Hall (1998) is an extension of Pearson's r to subset methods. It measures the linear correlation between each pair of selected features in addition to the correlation between selected features and the output class. The other subset algorithm to which we compared our work is MRMR Peng, Long, and Ding (2005). MRMR is an information-theory-based method that not only maximises the relevance of selected features to the supervised output class, but also minimises the redundancy among selected features. Different search strategies can be applied using these two algorithms as heuristics.

Although existing algorithms give various solutions, they are mostly ad-hoc, until Brown et al. (2012). Brown et al. derived an information-theoretical model from first principles, and retro-fitted existing information-theory-based feature selection methods to the model. In this paper, we consider the other side of the same coin. We proposed a probability-theory-based graphical model that has intuitive qualitative meaning in precision and recall, and we built a generalized feature selection scheme upon it.

A Probabilistic Model of Feature Selection

Feature selection is the selection of a subset of features F^* from the complete feature space F, for a specific classification/prediction problem, without loosing too much performance on some given measures. The benefit of feature selection include more explainable model, reduced computation time/power, and the avoidance of overfitting.

Existing feature selection algorithms, whose performance show no statistical difference on out benchmark, are mostly
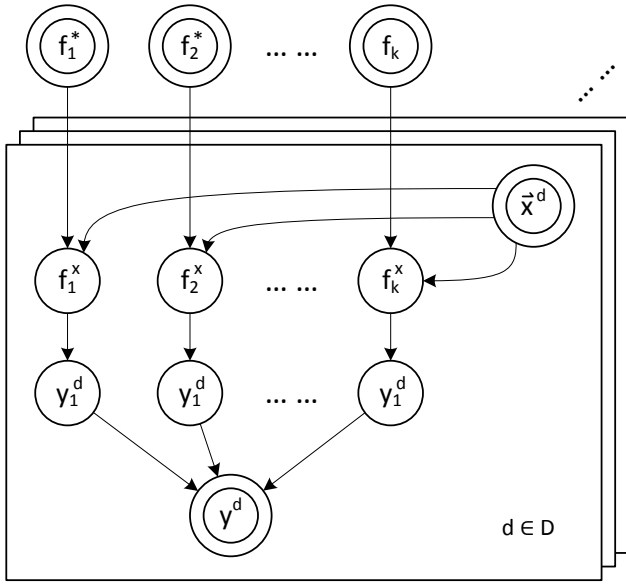
Figure 1: Graphical model. Double-wall circle denotes observed nodes, single-wall circle denotes hidden nodes.

$$f_k^* = \arg\max_{f_k} E_D \left[P \left(y^d | \vec{x}^d, F_k \right) \right]$$

$$= \arg\max_{f_k} \frac{1}{|D|} \sum_{d \in D} P \left(y^d | \vec{x}^d, F_k \right)$$

$$= \arg\max_{f_k} \sum_{d \in D} \frac{P \left(y^d, \vec{x}^d, F_k \right)}{P \left(\vec{x}^d, F_k \right)}$$

$$= \arg\max_{f_k} \sum_{d \in D} \sum_{\{y_i^d\}_{1 \le i \le k}} \frac{P \left(y^d, \vec{x}^d, F_k, \{y_i^d\}_{1 \le i \le k} \right)}{P \left(\vec{x}^d, F_k \right)}$$

$$= \arg\max_{f_k} \sum_{d \in D} \sum_{\{y_i^d\}_{1 \le i \le k}} \frac{P \left(y^d | \{y_i^d\}_{1 \le i \le k} \right) \left(\prod_{i=1}^{k} P \left(y_i^d | \vec{x}^d, f_i \right) \right) P (\tilde{x}}{P \left(\vec{x}^d, F_k \right)}$$

$$= \arg\max_{f_k} \sum_{d \in D} \sum_{\{y_i^d\}_{1 \le i \le k}} P \left(y^d | \{y_i^d\}_{1 \le i \le k} \right) \prod_{i=1}^{k} P \left(y_i^d | \vec{x}^d, f_i \right)$$
(1)

Here, we rewrote the expectation of a binary event as its probability, factorized the conditional probability in the joint probability divided by the marginal, marginalised over $\{y_i^d\}_{1 \le i \le k}$, factorized joint probability in conditional and prior following the graphical model and exploited d-separation to remove irrelevant conditions and to cancel terms in the equation.

Now we can optimise our initial objective aiming three goals of classification tasks: Precision, Recall, and General n-out-of-k.

Precision

In order to select the subset of features providing high precision binary classifier we need the agreement of all features predictors y_i^d to predict true. That is, when y^d is true we need all of the predictors y_i^d equals to true, and if y^d were false, just one of the predictors y_i^d would have to be false. Then the probability $P \left(y^d | \{y_i^d\}_{1 \le i \le k} \right)$ can be expressed as follows:

$$P \left(y^d | \{y_i^d\}_{1 \le i \le k} \right) = I \left[y^d = \bigwedge_{i=1}^{k} y_i^d \right] = \begin{cases} y^d = 1 \text{ when } \{y_i^d = 1\}_{1 \le i \le k} \\ y^d = 0 \text{ when } \{y_i^d\}_{1 \le i \le k, \exists y_j^d \|} \end{cases}$$
(2)

Then, we combined equations (1) and (2), separated each term according to the actual label value y^d and used the

---

ad-hoc. They give various answers to optimisation, but none of them gives the question. In this project, we started from the first principles, proposed a graphical model, and built a new feature selection system upon it.

Graphical Model

To begin the derivation, we propose a graphical model of feature selection for binary classification problem, as shown in figure 1. Double-circle nodes represent observed variables, including the data point vector $\vec{x}^d$, selected features f_i^* (where $1 \le i \le k$, $f_i \in F$) and the supervised class label y^d. Single-circle nodes represent latent variables, including a selected attribute of data point vector f_i^x, and prediction from a feature y_i^d.

Since jointly optimising this objective is NP-hard, we will build F^* in a greedy manner by choosing the next optimal feature f_k^* given the previous set of optimal features F_{k-1}^* and recursively defining $F_k^* = F_{k-1}^* \cup f_k^*$ with $F_0^* = \emptyset$. The measure of the fitness of a model can either be Expectation or Likelihood.

Expectation

The initial objective is defined as:

$$f_k^* = \arg\max_{f_k} E_D \left[P \left(y^d | \vec{x}^d, F_k \right) \right]$$

probability sum rule to rewire the second term as follows:

$$f_k^* = \arg\max_{f_k} \sum_{d \in D} I\left[y^d = \bigwedge_{i=1}^k y_i^d\right] \prod_{i=1}^k P\left(y_i^d | \vec{x}^d, f_i\right)$$

$$= \arg\max_{f_k} \sum_{d \in D} \begin{cases} y^d = 1 : \prod_{i=1}^k P\left(y_i^d = 1 | \vec{x}^d, f_i\right) \\ y^d = 0 : \sum_{\substack{\{y_i^d\}_{1 \le i \le k}, \\ \exists y_i^d \| y_i^d = 0}} \prod_{i=1}^k P\left(y_i^d | \vec{x}^d, f_i\right) \end{cases}$$

$$= \arg\max_{f_k} \sum_{d \in D} \begin{cases} y^d = 1 : \prod_{i=1}^k P\left(y_i^d = 1 | \vec{x}^d, f_i\right) \\ y^d = 0 : 1 - \prod_{i=1}^k P\left(y_i^d = 1 | \vec{x}^d, f_i\right) \end{cases}$$

$$(3)$$

$$= \arg\max_{f_k} \sum_{d \in D} \begin{cases} y^d = 1 : \prod_{i=1}^k P\left(y_i^d = 1 | \vec{x}^d, f_i\right) \\ y^d = 0 : 1 - \prod_{i=1}^k P\left(y^d = 1 | \vec{x}_i^d, f_i\right) \end{cases}$$

From (3) we can intuitively describe how this equation is related to the precision metric. In the confusion matrix, Precision is calculated dividing the numbers of true positives by the number of all examples predicted as positive (true positives and false positives). When y^d is true, (3) gives higher score to feature that have higher probability to predict true, encouraging true positives. Meanwhile, if y^d is false, (3) gives lower score to features that have higher probability to predict true, penalising false positives.

Recall case

In order to select the subset of features providing high recall binary classifier we need at least one y_i^d to predict true. Therefore, we need a disjunction operation between the predictors y_i^d. That is, when y^d is true, we need at least one predictor y_i^d to equals to true, and if y^d were false, all of the predictors y_i^d would have to be false. Then the probability $P\left(y^d | \{y_i^d\}_{1 \le i \le k}\right)$ can be expressed as follows:

$$P\left(y^d | \{y_i^d\}_{1 \le i \le k}\right) = I\left[y^d = \bigvee_{i=1}^k y_i^d\right] = \begin{cases} y^d = 1 \text{ when } \{y_i^d\}_{1 \le i \le k, \exists y_i^d \| y_i^d = 1} \\ y^d = 0 \text{ when } \{y_i^d = 0\}_{1 \le i \le k} \end{cases}$$

$$(4)$$

Here, we combined equations (1) and (4), separated each term according to the actual label value y^d and used the probability sum rule to rewrite the second term as follows:

$$f_k^* = \arg\max_{f_k} \sum_{d \in D} I\left[y^d = \bigvee_{i=1}^k y_i^d\right] \prod_{i=1}^k P\left(y_i^d | \vec{x}^d, f_i\right)$$

$$= \arg\max_{f_k} \sum_{d \in D} \begin{cases} y^d = 0 : \prod_{i=1}^k P\left(y_i^d = 0 | \vec{x}^d, f_i\right) \\ y^d = 1 : \sum_{\substack{\{y_i^d\}_{1 \le i \le k}, \\ \exists y_i^d \| y_i^d = 1}} \prod_{i=1}^k P\left(y_i^d | \vec{x}^d, f_i\right) \end{cases}$$

$$= \arg\max_{f_k} \sum_{d \in D} \begin{cases} y^d = 0 : \prod_{i=1}^k P\left(y_i^d = 0 | \vec{x}^d, f_i\right) \\ y^d = 1 : 1 - \prod_{i=1}^k P\left(y_i^d = 0 | \vec{x}^d, f_i\right) \end{cases}$$

$$(5)$$

$$= \arg\max_{f_k} \sum_{d \in D} \begin{cases} y^d = 0 : \prod_{i=1}^k P\left(y_i^d = 0 | \vec{x}^d, f_i\right) \\ y^d = 1 : 1 - \prod_{i=1}^k P\left(y^d = 0 | \vec{x}_i^d, f_i\right) \end{cases}$$

From (5) we can intuitively describe how this equation is related to the precision metric. In the confusion matrix, Recall is calculated dividing the numbers of actual true positives by the number of all true positives (true positives and false negatives). When y^d is false, (5) gives higher score to feature that have higher probability to predict false, encouraging true negatives. Meanwhile, if y^d is true, (5) gives lower score to features that have higher probability to predict false, penalizing false negatives.

Likelihood

$$f_k^* \equiv \arg\max_{f_k} \log \prod_{d \in D} P\left(y^d | \vec{x}^d, F_k\right)$$

$$= \arg\max_{f_k} \sum_{d \in D} \log \left(\frac{P\left(y^d, \vec{x}^d, F_k\right)}{P\left(\vec{x}^d, F_k\right)}\right)$$

$$= \arg\max_{f_k} \sum_{d \in D} \log \sum_{\{y_i^d\}_{1 \le i \le k}} \frac{P\left(y^d, \vec{x}^d, F_k, \{y_i^d\}_{1 \le i \le k}\right)}{P\left(\vec{x}^d, F_k\right)}$$

$$= \arg\max_{f_k} \sum_{d \in D} \log \sum_{\{y_i^d\}_{1 \le i \le k}} \frac{P\left(y^d | \{y_i^d\}_{1 \le i \le k}\right) \prod_{i=1}^k P\left(y_i^d | \vec{x}^d, f_i\right)}{P\left(\vec{x}^d, F_k\right)}$$

$$= \arg\max_{f_k} \sum_{d \in D} \log \sum_{\{y_i^d\}_{1 \le i \le k}} P\left(y^d | \{y_i^d\}_{1 \le i \le k}\right) \prod_{i=1}^k P\left(y_i^d | \vec{x}^d, f_i\right)$$

Here, we rewrote the likelihood of a binary event as its probability, factorized the conditional probability in the joint probability divided by the marginal, marginalized over $\{y_i^d\}_{1 \le i \le k}$, factorized joint probability in conditional and prior following the graphical model and exploited d-separation to remove irrelevant conditions and to cancel terms in the equation.

Since $\sum_{d \in D} \log \sum_{\{y_i^d\}_{1 \le i \le k}} P\left(y^d | \{y_i^d\}_{1 \le i \le k}\right) \prod_{i=1}^k P\left(y_i^d | \vec{x}^d, f_i\right)$ is exactly the same as the result in the expectation case, and the logarithm is a monotonous function, the choice of likelihood instead of expectation has no further affect on the algorithm.

Empirical

In this section, we conduct experiments to study the feature selection performances of several well known methods such as Symmetrical Uncertainty Rank (SUR), Gain Ratio Rank (GRR), Information Gain Rank (IGR), Correlation Based Rank (CBR), Conditional Entropy Rank (CER) Guyon and Elisseeff (2003), Correlation-based Feature Subset Selection (CFS) Hall (1998), Reliff (R) Robnik-Sikonja and Kononenko (2003), mRMR Peng, Long, and Ding (2005) and our proposed methods: High Recall Expectation (HRE) and High Recall Log likelihood (HRL), High Precision Expectation (HPE), High Precision Log likelihood (HPL). The first five methods are variable rank methods and the last seven are subset rank methods. The first five methods are variable rank methods that select variables by ranking them with some metric and the last seven are subset rank methods that assess subsets of variables and consider previous selections to decide the next selection.

Dataset	# Features	# Data	# Features/# Data	% True Labels	Feature Type
Breast Cancer	9	699	0.013	34 %	All Numeric
Diabetes	8	768	0.010	35 %	All Numeric
Heart Statlog	13	270	0.048	44 %	All Numeric
Spect	22	80	0.275	33 %	Categorical
Vote	16	435	0.037	39 %	Categorical
Newsgroup	500	1963	0.255	49 %	All Binary
Horse Colic	22	368	0.060	37 %	Categorical (16) and numeric (6)
Credit-American	15	690	0.022	44 %	Categorical(9) and numeric (6)
Credit-German	20	1000	0.020	30 %	Categorical (13) and Numeric (7)
Hepatitis	19	155	0.123	21 %	Categorical (13) and Numeric (6)
Ionosphere	34	351	0.097	36 %	All Numeric

Table 1: Statistics of the various datasets evaluated in this work.

TODO Rodrigo: please update acronyms in plots with those given above, use MRMR for MRMR... sorry, does not seem sensible to shorten this to 3 letters.

In this experimental study, we evaluated each feature selection method on a variety of binary classification problems from the UCI machine learning repository Bache and Lichman (2013). These datasets along with their properties are outlined in Table 1.

In addition, we used three different classifiers to solve each of the binary classifier problem. Logistic Regression (LR), SVM Linear (SVM) and Naive Bayes (NB) were used in the experiments and the first two are implemented in the LibLinear library Fan et al. (2008) and the last one is implemented in the Weka software Hall et al. (2009).

To evaluate feature selection algorithm performance, we perform 10-fold nested cross-validation (nesting for tuning hyperparameters of each classification algorithm) and evaluate accuracy, precision, and recall at each stage of feature selection from the first selected feature through to the maximum (either 50 or the number of features in the dataset, whichever is smaller).

In our evaluation, we wish to answer the following questions: Do different feature selection methods perform better with each classifier or on each dataset in comparison to others? How reliably do the different feature selection methods perform overall in terms of their performance distribution?

To answer these questions, we first start with Figure 2, where we evaluate the performance of various feature selection algorithms per dataset, averaged across classifier type and each stage of feature selection. We observe...

In Figure 3, we examine how well our novel feature selection algorithms HRE, HRL, HPE, HPL perform vs. classifier, noting that overall performance is best for Naive Bayes followed by Logistic Regression and noticeably worse for the SVM — it seems simply that the probabilistic nature of our feature selection dovetail best with classification methods that are themselves probabilistic, and overall best with Naive Bayes which makes the same independence assumptions. **TODO Rodrigo: can we get bargraphs like those in FeatureSelection classifier, except only showing HRE, HRL, HPE, HPL *and* averaging over all datasets? So major group=HRE, HRL, HPE, HPL, minor group=SVM,LR,NB and one graph for each**

of A/P/R/F. Also the filename should not have spaces.

Next we examine Figure 4, where we evaluate the performance distribution of various feature selection algorithms with samples taken for each dataset, classifier type, and stage of feature selection. Here we see Y. We observe...

Conclusion

References

Bache, K., and Lichman, M. 2013. UCI machine learning repository.

Brown, G.; Pocock, A.; Zhao, M.-J.; and Luján, M. 2012. Conditional likelihood maximisation: A unifying framework for information theoretic feature selection. *The Journal of Machine Learning Research* 13:27–66.

Fan, R.-E.; Chang, K.-W.; Hsieh, C.-J.; Wang, X.-R.; and Lin, C.-J. 2008. LIBLINEAR: A library for large linear classification. *Journal of Machine Learning Research* 9:1871–1874.

Guyon, I., and Elisseeff, A. 2003. An introduction to variable and feature selection. *J. Mach. Learn. Res.* 3:1157–1182.

Guyon, I.; Weston, J.; Barnhill, S.; and Vapnik, V. 2002. Gene selection for cancer classification using support vector machines. *Machine learning* 46(1-3):389–422.

Hall, M.; Frank, E.; Holmes, G.; Pfahringer, B.; Reutemann, P.; and Witten, I. H. 2009. The weka data mining software: An update. *SIGKDD Explorations* 11.

Hall, M. A. 1998. *Correlation-based Feature Subset Selection for Machine Learning*. Ph.D. Dissertation, University of Waikato, Hamilton, New Zealand.

Kira, K., and Rendell, L. A. 1992. The feature selection problem: Traditional methods and a new algorithm. In *AAAI*, 129–134.

Peng, H.; Long, F.; and Ding, C. 2005. Feature selection based on mutual information: Criteria of max-dependency, max-relevance, and min-redundancy. *Pattern Analysis and Machine Intelligence, IEEE Transactions* 27:1226–1238.

Figure 2: Performance of various feature selection algorithms per dataset, averaged across classifier type and each stage of feature selection.


Figure 3: Performance of various feature selection algorithms per classifier, averaged across dataset and each stage of feature selection.
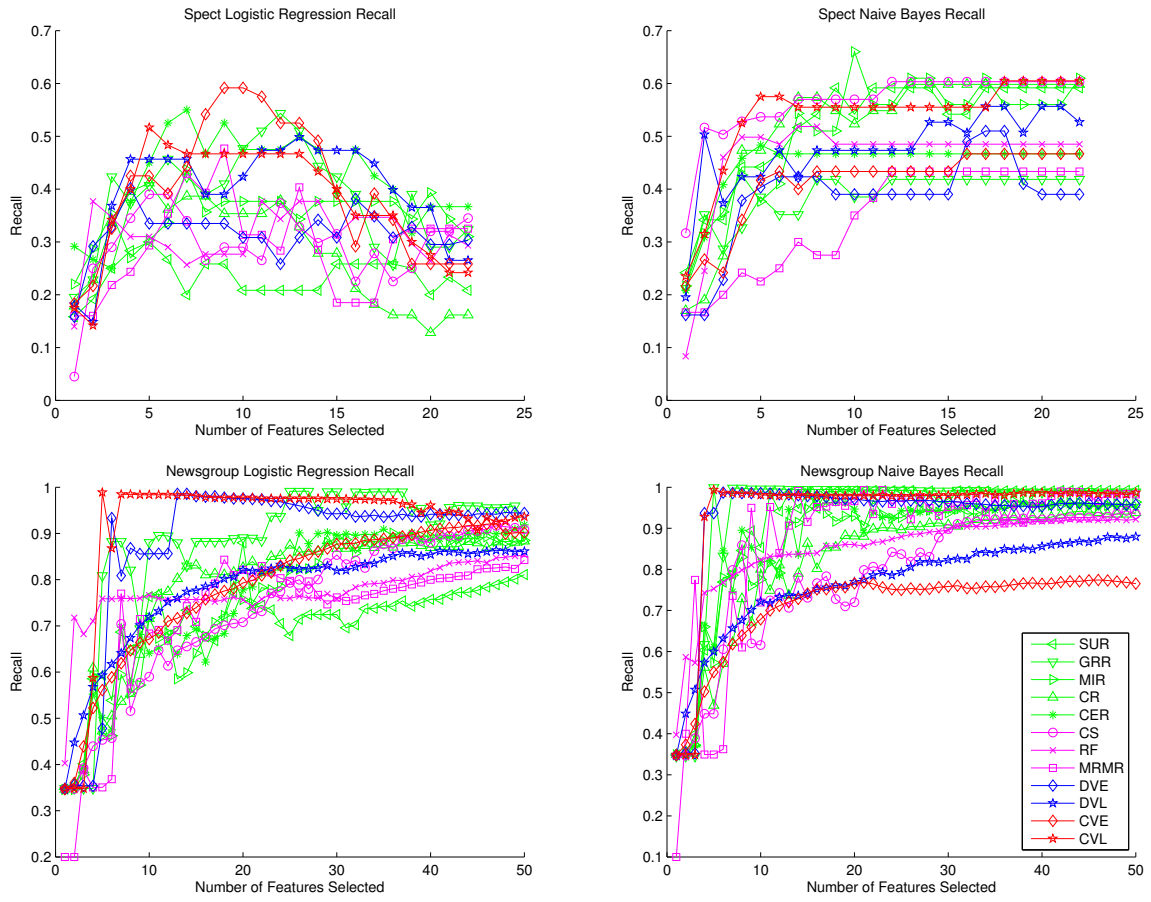
Figure 4: Performance distribution of various feature selection algorithms with samples taken for each dataset, classifier type, and stage of feature selection.

Robnik-Sikonja, M., and Kononenko, I. 2003. Theoretical and empirical analysis of relieff and rrelieff. *Machine Learning* 53:23–69.